

OpenACC

More Science, Less Programming

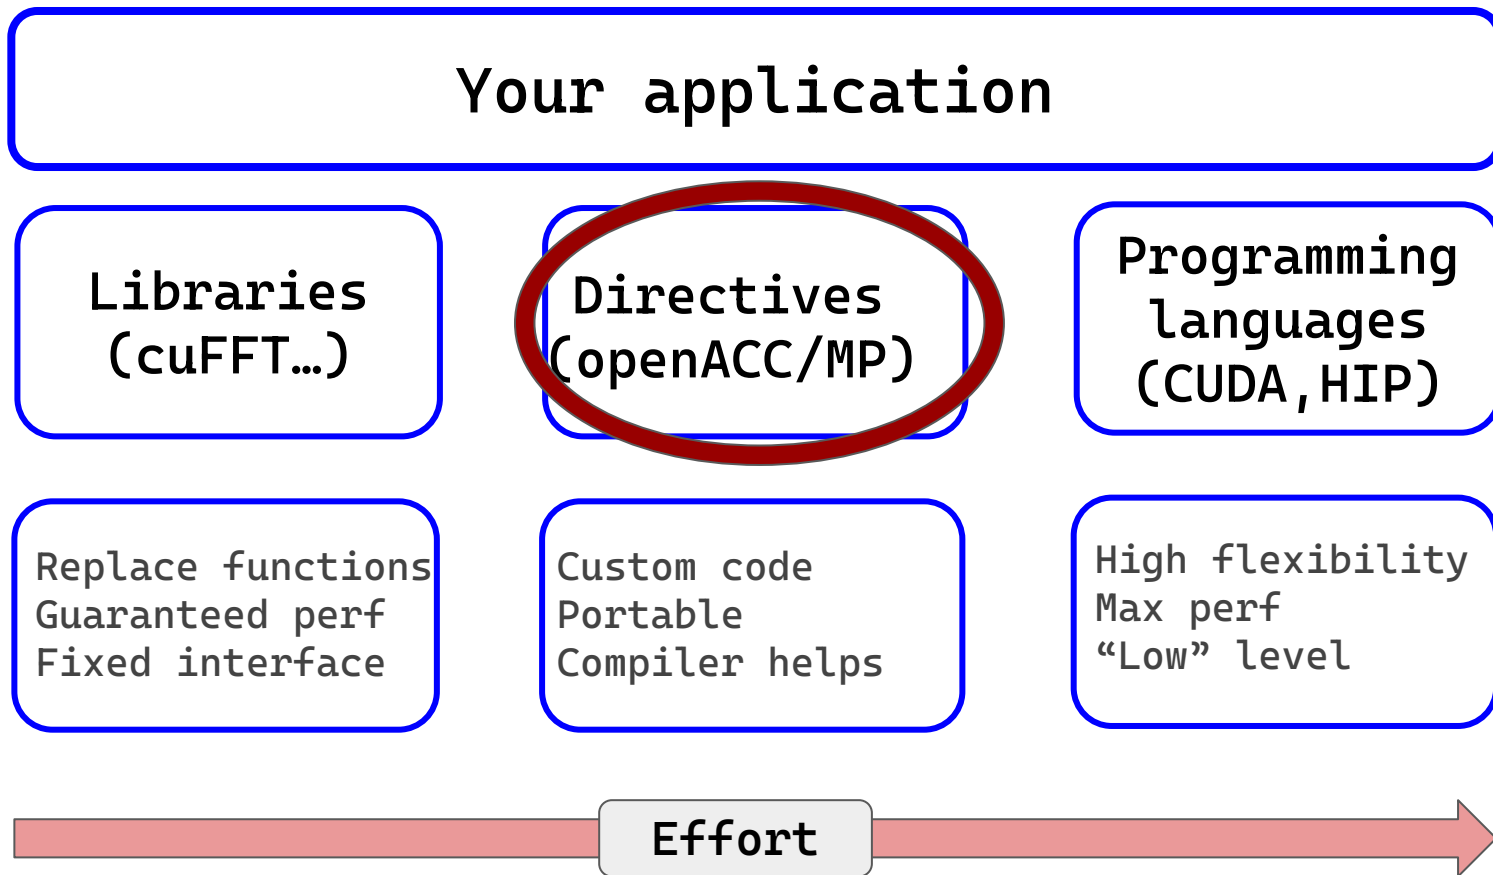
First steps in OpenACC

Francesco Andreucci
SISSA

ICTP CIFRA Magurele School,
3/07/2025



How write a code that exploits a GPU



How write a code that exploits a GPU

Your application

<https://www.openacc.org/resources>

OpenACC

More Science, Less Programming

R
H
Fixed Interface

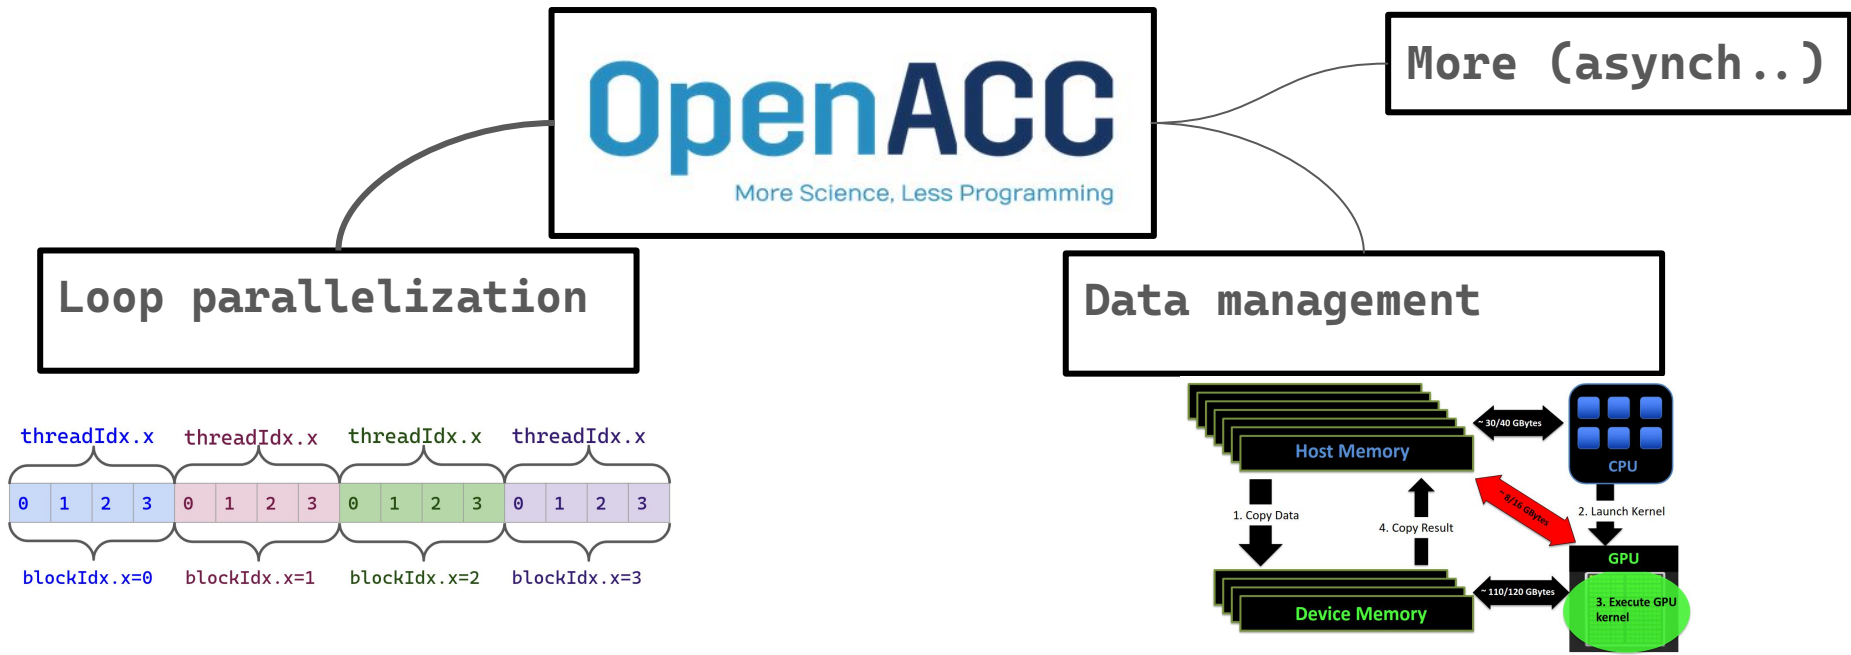
Compiler helps

Low level

Effort


What can you do with OpenACC?

- Runs on different compilers (nvc/pgi, gcc...) and different accelerators (NVIDIA,AMD,multi core CPU...)



First taste of OpenACC

```
for (int i=0; i<N; i++){  
    a[i]=1; b[i]=2;  
}
```



```
#pragma acc parallel loop  
for (int i=0; i<N; i++){  
    a[i]=1; b[i]=2;  
}
```


The compiler interprets your directives and **generates GPU CUDA code**

OpenACC is directive-based:

- Directives are **incrementally** added on top of the serial cpu-code (**no refactoring/code duplication**)
- If the compiler does not recognize a directive it **is treated as a comment**
- Supports both C/C++ and Fortran
- Handles both **loop parallelization** and **data management**

First taste of OpenACC

```
do i=1,N  
  a(i)=1  
end do
```



```
!$acc parallel loop  
do i=1,N  
  a(i)=1;  
end do
```

The compiler interprets your directives and **generates GPU CUDA code**

OpenACC is directive-based:

- Directives are **incrementally** added on top of the serial cpu-code (**no refactoring/code duplication**)
- If the compiler does not recognize a directive it **is treated as a comment**
- Supports both C/C++ and Fortran
- Handles both **loop parallelization** and **data management**

Pragma-directives syntax

C/C++	<code>#pragma acc {directives} {clauses}</code>
Fortran	<code>!\$acc {directives} {clauses}</code>

- **Directives** are **standalone** commands that tell the compiler to alter the code in some way (i.e memory copies, parallelize loops...)
- **Clauses** are **auxiliary** specifiers that refine the action of the directives

How to compile(on Leonardo, Ampere A100 GPU cc=80):

```
nvc          -o exe myprog.c      -acc -gpu=cc80,cuda12.3 -Minfo=acc
nvc++        -o exe myprog.cpp    -acc -gpu=cc80,cuda12.3 -Minfo=acc
nvfortran    -o exe myprog.f90    -acc -gpu=cc80,cuda12.3 -Minfo=acc
```



Let's move on the cluster..

-Minfo=all output

```
#pragma acc parallel loop
for(int i=0; i<n; i++){
    a[i]=3;
    b[i]=4;
}
```

```
#pragma acc parallel loop
for(int j=0; j<n; j++){
    a[j] = 2*a[j];
    b[j] = b[j]+1;
}
```

```
#pragma acc parallel loop
for(int k=0; k<n; k++){
    c[k] = a[k]+b[k];
}
```

-Minfo=all output

```
#pragma omp main:
for(i=0; i<n; i++)
{
    12, Generating implicit firstprivate(i,n)
    Generating NVIDIA GPU code
    15, #pragma acc loop gang, vector(128) /* blockIdx.x threadIdx.x */
    12, Generating implicit copyout(b[:100000],a[:100000]) [if not already present]
    18, Generating implicit firstprivate(n,j)
    Generating NVIDIA GPU code
    21, #pragma acc loop gang, vector(128) /* blockIdx.x threadIdx.x */
    18, Generating implicit allocate(a[:100001]) [if not already present]
    Generating implicit copyin(a[1:100000]) [if not already present]
    Generating implicit copy(b[:100000]) [if not already present]
    Generating implicit copyout(a[:100000]) [if not already present]
    24, Generating implicit firstprivate(n,k)
    Generating NVIDIA GPU code
    27, #pragma acc loop gang, vector(128) /* blockIdx.x threadIdx.x */
    24, Generating implicit copyin(a[:100000]) [if not already present]
    Generating implicit copyout(c[:100000]) [if not already present]
    Generating implicit copyin(b[:100000]) [if not already present]
}
```

-Minfo=all output

```
#pragma acc parallel loop
for(int i=0; i<n; i++){
    a[i]=3;
    b[i]=4;
}
```

```
#pragma acc parallel loop
for(int j=0; j<n; j++){
    a[j] = 2*a[j];
    b[j] = b[j]+1;
}
```

```
#pragma acc parallel loop
for(int k=0; k<n; k++){
    c[k] = a[k]+b[k];
}
```

```
Generating implicit firstprivate(i,N)
Generating NVIDIA GPU code
22, #pragma acc loop gang, vector(128)
    /* blockIdx.x threadIdx.x */
Generating implicit copyout(b[:10000000],a[:10000000])
[if not already present]
```

-Minfo=all output

```
#pragma acc parallel loop
for(int i=0; i<n; i++){
    a[i]=3;
    b[i]=4;
}
```

```
Generating implicit firstprivate(i,N)
Generating NVIDIA GPU code
22, #pragma acc loop gang, vector(128)
    /* blockIdx.x threadIdx.x */
Generating implicit copyout(b[:10000000],a[:10000000])
[if not already present]
```

```
#pragma acc parallel loop
for(int j=0; j<n; j++){
    a[j] = 2*a[j];
    b[j] = b[j]+1;
}
```

```
21, #pragma acc loop gang, vector(128) /
Generating implicit allocate(a[:100001])
Generating implicit copyin(a[1:100000])
Generating implicit copy(b[:100000]) [if
Generating implicit copyout(a[:100000])
Generating implicit firstprivate(n,k)
Generating NVIDIA GPU code
```

```
#pragma acc parallel loop
for(int k=0; k<n; k++){
    c[k] = a[k]+b[k];
}
```

-Minfo=all output

```
#pragma acc parallel loop
for(int i=0; i<n; i++){
    a[i]=3;
    b[i]=4;
}
```

```
Generating implicit firstprivate(i,N)
Generating NVIDIA GPU code
22, #pragma acc loop gang, vector(128)
    /* blockIdx.x threadIdx.x */
Generating implicit copyout(b[:10000000],a[:10000000])
[if not already present]
```

```
#pragma acc parallel loop
for(int j=0; j<n; j++){
    a[j] = 2*a[j];
    b[j] = b[j]+1;
}
```

```
21, #pragma acc loop gang, vector(128) /
Generating implicit allocate(a[:100001])
Generating implicit copyin(a[1:100000])
Generating implicit copy(b[:100000]) [if
Generating implicit copyout(a[:100000])
Generating implicit firstprivate(n,k)
Generating NVIDIA GPU code
```

```
#pragma acc parallel loop
for(int k=0; k<n; k++){
    c[k] = a[k]+b[k];
}
```

```
Generating implicit firstprivate(j,N)
Generating NVIDIA GPU code
34, #pragma acc loop gang, vector(128)
    /* blockIdx.x threadIdx.x */
Generating implicit copyin(a[:10000000])
Generating implicit copyout(c[:10000000])
Generating implicit copyin(b[:10000000])
```


Memory transfers

```
#pragma acc parallel loop
for(int i=0; i<n; i++){
    a[i]=3;
    b[i]=4;
}
```

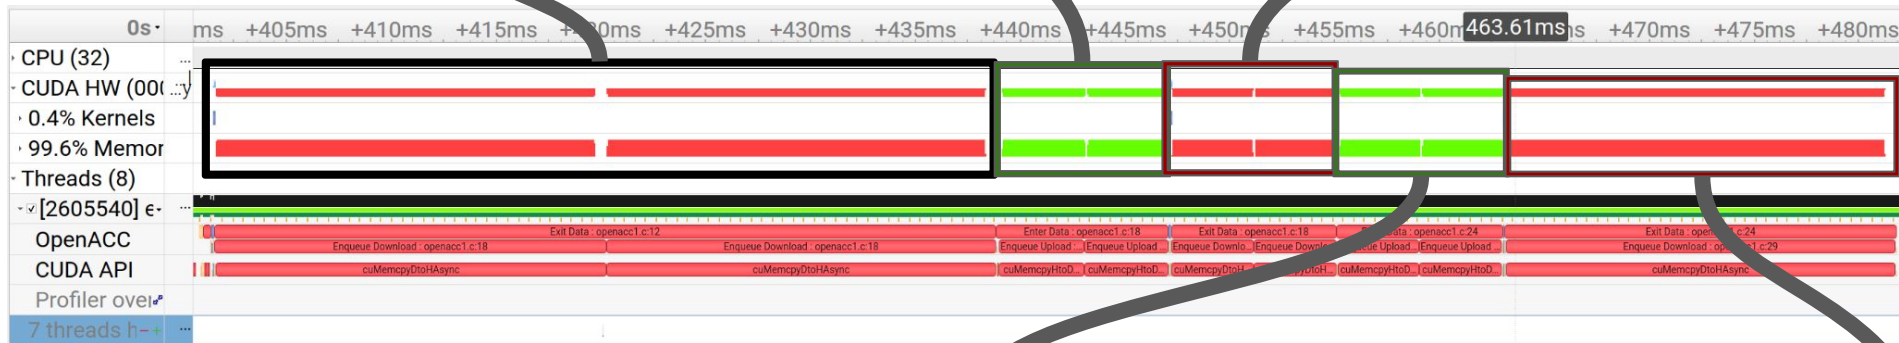
out a,b

out a, b

in a, b

```
#pragma acc parallel loop
for(int j=0; j<n; j++){
    a[j] = 2*a[j];
    b[j] = b[j]+1;
}
```

out a, b



in a, b

```
#pragma acc parallel loop
for(int k=0; k<n; k++){
    c[k] = a[k]+b[k];
}
```

```
out c
```

Memory transfers

```
#pragma acc parallel loop
for(int i=0; i<n; i++){
    a[i]=3;
    b[i]=4;
}
```

out a,b

in a,b

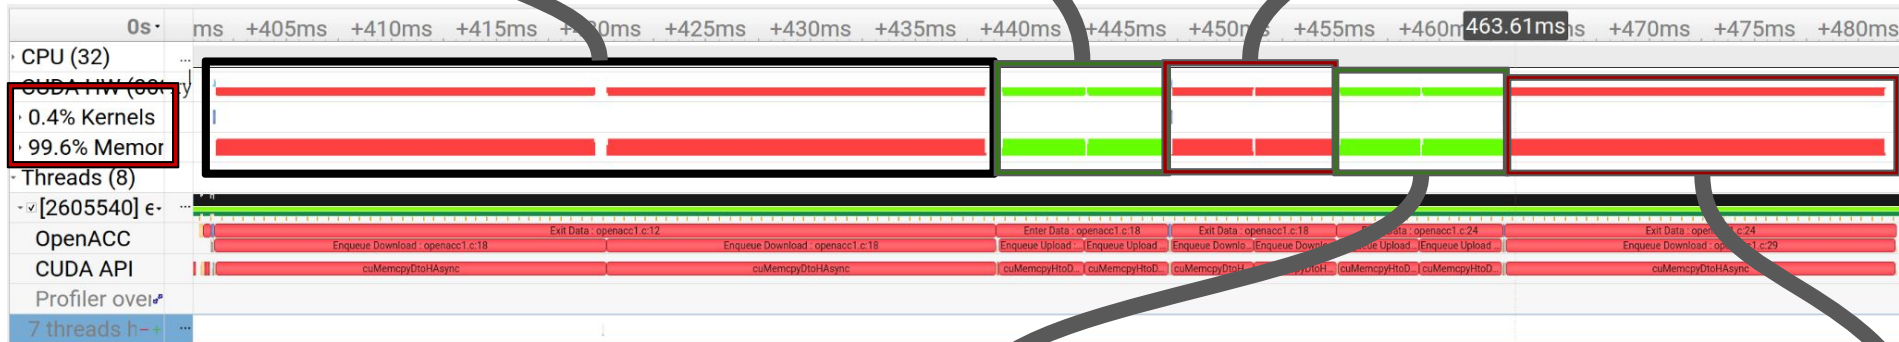
```
#pragma acc parallel loop
for(int j=0; j<n; j++){
    a[j] = 2*a[j];
    b[j] = b[j]+1;
}
```

out a,b

in a,b

```
#pragma acc parallel loop
for(int k=0; k<n; k++){
    c[k] = a[k]+b[k];
}
```

out c



...99.6% of the GPU activity is spent in memory transfers!

Data Management

```
#pragma acc parallel loop
for(int i=0; i<n; i++){
    a[i]=3;
    b[i]=4;
}
```

```
#pragma acc parallel loop
for(int j=0; j<n; j++){
    a[j] = 2*a[j];
    b[j] = b[j]+1;
}
```

```
#pragma acc parallel loop
for(int k=0; k<n; k++){
    c[k] = a[k]+b[k];
}
```

- What we just saw is called **implicit** data management
- This can be **very inefficient** especially if the parallel directive is iteratively invoked inside a loop

Data Management: main points

- Data must be visible to the device when running in a parallel region
- After a parallel region, we might want to get back the updated data from the device on the host
- To maximize the performance one should minimize data transfers between host/device

Data clauses

- Data clauses are used to tell the compiler which data we want to move

```
#pragma acc parallel loop copy(a[0:n])  
for(i=0;i<n;i++){  
    a[i] = a[i] +1;  
}
```

a[start_index:length]

C/++

a(start_index:end_index)

F90

- You can also slice the array and copy only what you need

```
#pragma acc parallel loop copy(a[5:m])  
for(i=5;i<m;i++){  
    a[i] = a[i] +1;  
}
```

Data clauses

- Data clauses are used to tell the compiler which data we want to move

```
!$acc parallel loop copy(a(1:n))  
do i=1,n  
    a(i) = a(i) +1  
end do
```

a[start_index:length]

C/++

a(start_index:end_index)

F90

- You can also slice the array and copy only what you need

```
!$acc parallel loop copy(a(5:m))  
do i=5,m  
    a(i) = a(i) +1  
end do
```

Data clauses

- | | |
|--|---|
| <ul style="list-style-type: none">• <code>copy(list)</code> | <ul style="list-style-type: none">• Allocates memory on GPU, copies data from host when entering the region and copies data back from device at the end |
| <ul style="list-style-type: none">• <code>copyin(list)</code> | <ul style="list-style-type: none">• Allocates memory on GPU, copies data from host when entering the region |
| <ul style="list-style-type: none">• <code>copyout(list)</code> | <ul style="list-style-type: none">• Allocates memory on GPU, copies data back from device at the end |
| <ul style="list-style-type: none">• <code>create(list)</code> | <ul style="list-style-type: none">• Allocates memory on GPU (useful for tmp buffers) |
| <ul style="list-style-type: none">• <code>delete(list)</code> | <ul style="list-style-type: none">• Frees memory on the GPU(!) |

Structured data regions

```
#pragma acc parallel loop
for(int i = 0; i < N; i++){
    a[i] = 3; b[i]=4;
}

#pragma acc parallel loop
for(int i = 0; i < N; i++){
    a[i] = 2*a[i]; b[i] = b[i]+1
}

#pragma acc parallel loop
for(int i = 0; i < N; i++){
    c[i] = a[i] + b[i];
}
```

a copied in and out

a,b copied in and out

a,b copied in and c copied
out

Structured data regions

```
#pragma acc data copyin(a[0:N],b[0:N])
copyout(c[0:N])
{
    #pragma acc parallel loop
    for(int i = 0; i < N; i++){
        a[i] = 3; b[i]=4;
    }
    #pragma acc parallel loop
    for(int i = 0; i < N; i++){
        a[i] = 2*a[i]; b[i] = b[i]+1
    }
    #pragma acc parallel loop
    for(int i = 0; i < N; i++){
        c[i] = a[i] + b[i];
    }
}
```

- Inside the **data region** the runtime knows a,b,c are in the GPU and does not copy them back and forth
- As we saw before, a lone **parallel(/kernels)** directive open an **implicit data region** that spans the parallel region

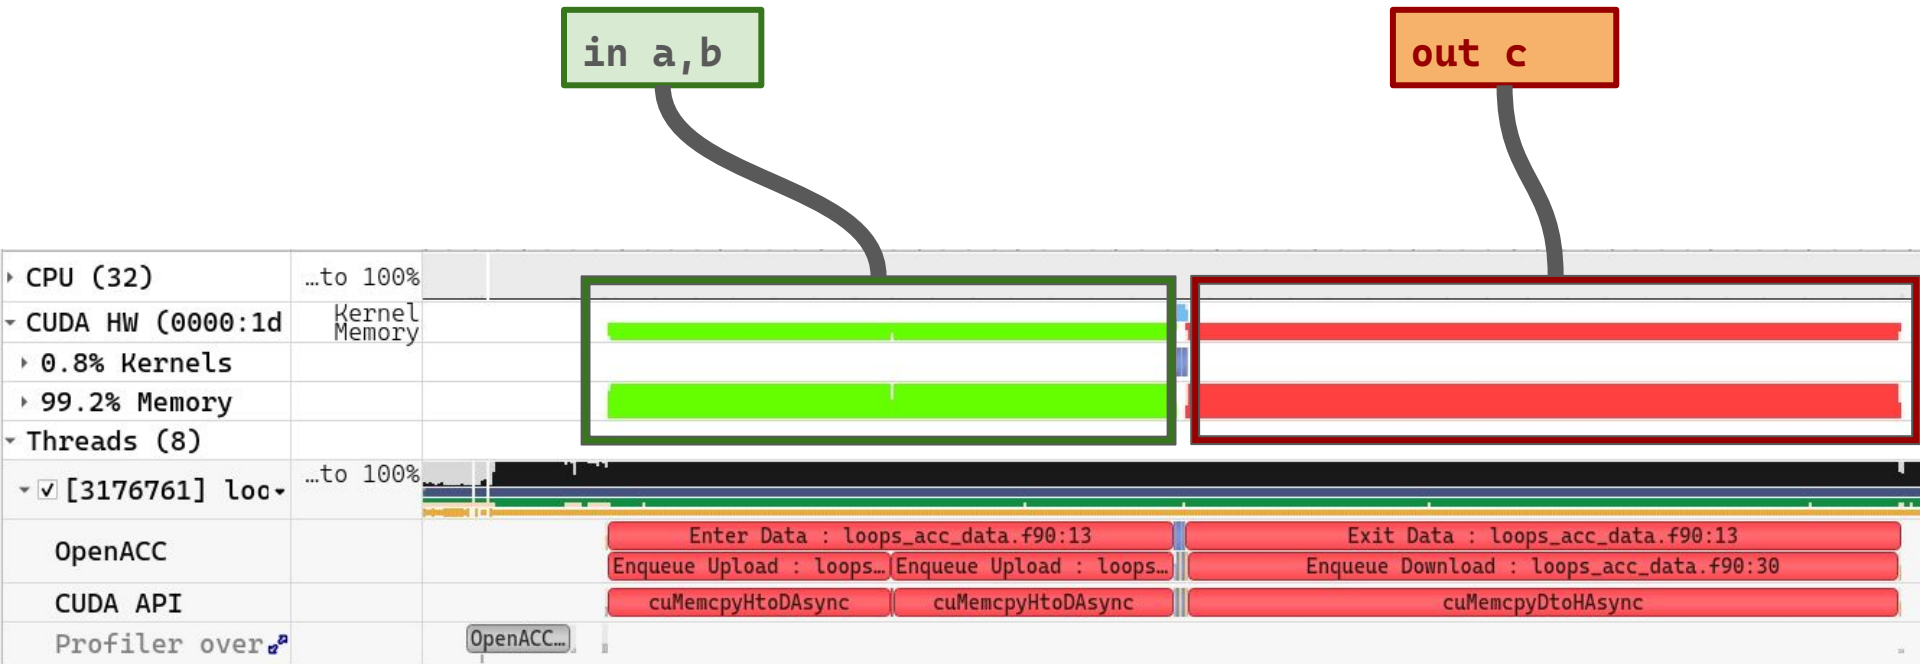
Structured data regions

```
#pragma acc data copyin(a[0:N],b[0:N])
copyout(c[0:N])
{
    #pragma acc parallel loop
    for(int i = 0; i < N; i++){
        a[i] = 3; b[i]=4;
    }
    #pragma acc parallel loop
    for(int i = 0; i < N; i++){
        a[i] = 2*a[i]; b[i] = b[i]+1
    }
    #pragma acc parallel loop
    for(int i = 0; i < N; i++){
        c[i] = a[i] + b[i];
    }
}
```

main:

```
15, Generating copyin(a[:n],b[:n])
    [if not already present]
    Generating copyout(c[:n])
    [if not already present]
    Generating implicit firstprivate(n,i)
    Generating NVIDIA GPU code
17, #pragma acc loop gang, vector(128)
    /* blockIdx.x threadIdx.x */
20, Generating implicit firstprivate(n,j)
    Generating NVIDIA GPU code
31, #pragma acc loop gang, vector(128)
    /* blockIdx.x threadIdx.x */
34, Generating implicit firstprivate(n,k)
    Generating NVIDIA GPU code
37, #pragma acc loop gang, vector(128)
    /* blockIdx.x threadIdx.x */
```

Structured data regions



Unstructured data regions

```
#pragma acc enter data \ copyin(a[0:N],b[0:N])  
create(c[0:N])
```

```
#pragma acc parallel loop  
for(int i = 0; i < N; i++){  
    a[i] = 3; b[i]=4;  
}
```

```
#pragma acc parallel loop  
for(int i = 0; i < N; i++){  
    a[i] = 2*a[i]; b[i] = b[i]+1  
}
```

```
#pragma acc parallel loop  
for(int i = 0; i < N; i++){  
    c[i] = a[i] + b[i];  
}
```

```
#pragma acc exit data delete(a,b) \copyout(c)
```

- **enter data**
+ copyin, create
- **exit data**
+ copyout, delete

(Un)structured data regions

Unstructured data

- Multiple start/end points
- Can branch across scopes
- Memory must be explicitly deallocated

```
#pragma acc enter data \
    copyin(a[0:n],b[0:n]) create(c[0:n])

#pragma acc parallel loop
    for(i=0;i<n;i++) c[i]=a[i]+b[i];

#pragma acc exit data copyout(c[0:n])
delete(a,b)
```

Structured data

- Explicit start/end points
- Within a single scope
- Memory is deallocated after the data region

```
#pragma acc data copyin(a[0:n],b[0:n])\
    copyout(c[0:n])
{
    #pragma acc parallel loop
        for(i=0;i<n;i++) c[i]=a[i]+b[i];
}
```

Unstructured data regions

```
int * allocate(int size)
{
    int * ptr = (int *) malloc(size*sizeof(int));

    return ptr;
}

void deallocate(int * ptr)
{
    free(ptr);
}

int main()
{
    int *ptr = allocate(100);

    for(int i=0; i<100; i++){
        ptr[i]=0;
    }
    deallocate(ptr)

    return 0;
}
```

Unstructured data regions

Data allocation

```
int * allocate(int size)
{
    int * ptr = (int *) malloc(size*sizeof(int));
    return ptr;
}
```

```
void deallocate(int * ptr)
{
    free(ptr);
}
```

```
int main()
{
    int *ptr = allocate(100);
    for(int i=0; i<100; i++){
        ptr[i]=0;
    }
    deallocate(ptr)

    return 0;
}
```

Unstructured data regions

```
int * allocate(int size)
{
    int * ptr = (int *) malloc(size*sizeof(int));

    return ptr;
}

void deallocate(int * ptr)
{
    free(ptr);
}

int main()
{
    int *ptr = allocate(100);
    for(int i=0; i<100; i++){
        ptr[i]=0;
    }
    deallocate(ptr)

    return 0;
}
```

Data allocation

Data utilization

Unstructured data regions

```
int * allocate(int size)
{
    int * ptr = (int *) malloc(size*sizeof(int));

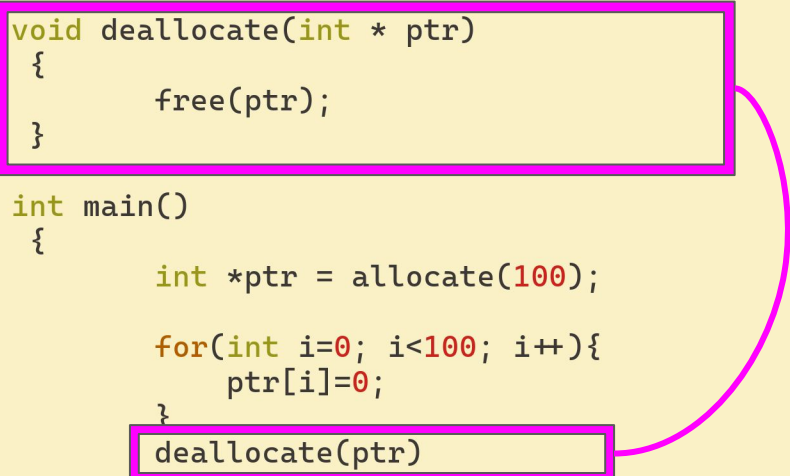
    return ptr;
}

void deallocate(int * ptr)
{
    free(ptr);
}

int main()
{
    int *ptr = allocate(100);

    for(int i=0; i<100; i++){
        ptr[i]=0;
    }
    deallocate(ptr)

    return 0;
}
```



Data allocation

Data utilization

Data deallocation

Unstructured data regions

```
int * allocate(int size)
{
    int * ptr = (int *) malloc(size*sizeof(int));
    #pragma acc enter data create(c[0,size])
    return ptr;
}

void deallocate(int * ptr)
{
    #pragma acc exit data delete(ptr)
    free(ptr);
}

int main()
{
    int *ptr = allocate(100);
    #pragma acc parallel loop
    for(int i=0; i<100; i++){
        ptr[i]=0;
    }
    deallocate(ptr)

    return 0;
}
```

Data allocation

Data utilization

Data deallocation

Unstructured data regions

```
int * allocate(int size)
{
    int * ptr = (int *) malloc(size*sizeof(int));
    #pragma acc enter data create(c[0,size])
    return ptr;
}

void deallocate(int * ptr)
{
    #pragma acc exit data delete(ptr)
    free(ptr);
}

int main()
{
    int *ptr = allocate(100);
    #pragma acc parallel loop
    for(int i=0; i<100; i++){
        ptr[i]=0;
    }
    deallocate(ptr)

    return 0;
}
```

Data allocation

Data deallocation

Data utilization

Keeping host and device in synch

```
#pragma acc data copyin(a[0:n],b[0:n]) \
    copyout(c[0:n])
{
    #pragma acc parallel loop
    for(int i=0;i<n;i++){
        a[i]=3; b[i]=4;
    }

    printf(a); printf(b);

    for(int j=0;j<n;j++){
        a[j]=2*a[j]; b[j]=b[j]+1;
    }
    #pragma acc parallel loop
    for(int k=0;k<n;k++){
        c[k] = a[k] + b[k];
    }
}

printf(c);
```

Keeping host and device in synch

```
#pragma acc data copyin(a[0:n],b[0:n]) \  
    copyout(c[0:n])
```

```
{
```

```
    #pragma acc parallel loop  
    for(int i=0;i<n;i++){  
        a[i]=3; b[i]=4;  
    }
```

```
    printf(a); printf(b);
```

```
    for(int j=0;j<n;j++){  
        a[j]=2*a[j]; b[j]=b[j]+1;  
    }
```

```
    #pragma acc parallel loop  
    for(int k=0;k<n;k++){  
        c[k] = a[k] + b[k];  
    }
```

```
}
```

```
    printf(c);
```

a and b are updated **on device**

Keeping host and device in synch

```
#pragma acc data copyin(a[0:n],b[0:n]) \  
    copyout(c[0:n])
```

```
{
```

```
    #pragma acc parallel loop  
    for(int i=0;i<n;i++){  
        a[i]=3; b[i]=4;  
    }
```

```
    printf(a); printf(b);
```

```
    for(int j=0;j<n;j++){  
        a[j]=2*a[j]; b[j]=b[j]+1;  
    }
```

```
    #pragma acc parallel loop  
    for(int k=0;k<n;k++){  
        c[k] = a[k] + b[k];  
    }
```

```
}
```

```
    printf(c);
```

a and b are updated **on device**

a and b are **printed from host (un updated)**

Keeping host and device in synch

```
#pragma acc data copyin(a[0:n],b[0:n]) \  
    copyout(c[0:n])
```

```
{
```

```
#pragma acc parallel loop  
for(int i=0;i<n;i++){  
    a[i]=3; b[i]=4;  
}
```

```
printf(a); printf(b);
```

```
for(int j=0;j<n;j++){  
    a[j]=2*a[j]; b[j]=b[j]+1;  
}
```

```
#pragma acc parallel loop  
for(int k=0;k<n;k++){  
    c[k] = a[k] + b[k];  
}
```

```
}
```

```
printf(c);
```

a and b are updated on **device**

a and b are **printed from host (un updated)**

a and b are **updated on host**

Keeping host and device in synch

```
#pragma acc data copyin(a[0:n],b[0:n]) \  
    copyout(c[0:n])
```

```
{
```

```
#pragma acc parallel loop  
for(int i=0;i<n;i++){  
    a[i]=3; b[i]=4;  
}
```

```
printf(a); printf(b);
```

```
for(int j=0;j<n;j++){  
    a[j]=2*a[j]; b[j]=b[j]+1;  
}
```

```
#pragma acc parallel loop  
for(int k=0;k<n;k++){  
    c[k] = a[k] + b[k];  
}
```

```
}
```

```
printf(c);
```

a and b are updated on **device**

a and b are **printed from host** (un updated)

a and b are **updated on host**

c is **updated on device**
c is **copied out to host**

Keeping host and device in synch

```
#pragma acc data copyin(a[0:n],b[0:n]) \  
    copyout(c[0:n])
```

```
{
```

```
#pragma acc parallel loop  
for(int i=0;i<n;i++){  
    a[i]=3; b[i]=4;  
}
```

```
printf(a); printf(b);
```

```
for(int j=0;j<n;j++){  
    a[j]=2*a[j]; b[j]=b[j]+1;  
}
```

```
#pragma acc parallel loop  
for(int k=0;k<n;k++){  
    c[k] = a[k] + b[k];  
}
```

```
}
```

```
printf(c);
```

a and b are updated on **device**

a and b are **printed from host** (un updated)

a and b are **updated on host**

c is **updated on device**
c is **copied out to host**

c is **printed from host**

Keeping host and device in synch

If you try to copy data already PRESENT on the GPU, the copy is not done.

< a[] modified on the host >

#pragma acc enter data copyin(a[0:N]) → host and device copies are in sync

< a modified on the host > → host and device copies are out of sync

#pragma acc data copyin(a[0:N]) ! copy is ignored

< a used on the GPU > → host and device copies are out of sync

Keeping host and device in synch

If you try to copy data already PRESENT on the GPU, the copy is not done.

<a[] is changed on device>

```
#pragma acc update host(a[0:n]) }
```

a is copied from GPU to CPU

<a[] is changed on host>

```
#pragma acc update device(a[0:n]) }
```

a is copied from CPU to GPU

- You can also shape the update: `update host(a[start_index:count])`
- You can also use: `update self = update host`

Keeping host and device in synch

If you try to copy data already PRESENT on the GPU, the copy is not done.

<a() is changed on device>

!\$acc update host(a(1:n))

}

a is copied from GPU to CPU

<a() is changed on host>

!\$acc update device(a(1:n))

}

a is copied from CPU to GPU

- You can also shape the update: `update host(a(start_index:end_index))`
- You can also use: `update self = update host`



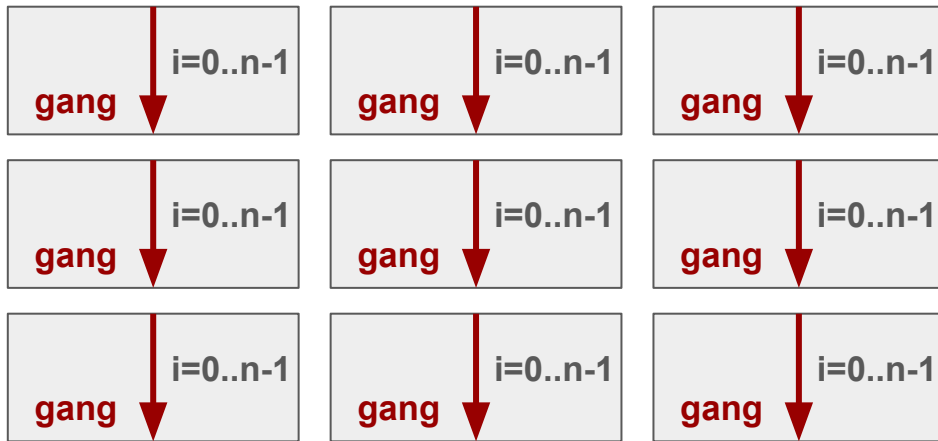
Let's move on the cluster..

The parallel directive

```
#pragma acc parallel
{
    for(i=0;i<n;i++){
        <code>
    }
}
```

The **parallel** directive:

- Creates a region where the code is run on the GPU **sequentially** (single thread)
- A **gang** is an **independent group of threads** (more details later...)

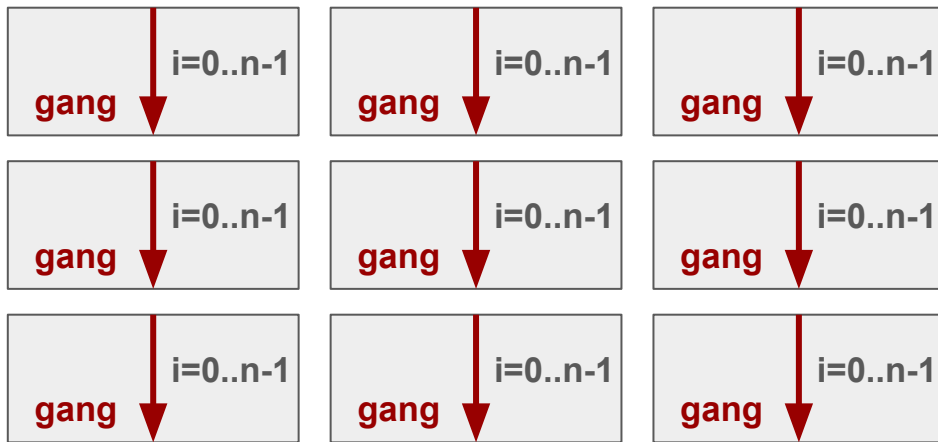


The parallel directive

```
#pragma acc parallel
{
    for(i=0;i<n;i++){
        <code>
    }
}
```

The **parallel** directive:

- Creates a region where the code is run on the GPU **sequentially** (single thread)
- A **gang** is an **independent group of threads** (more details later...)



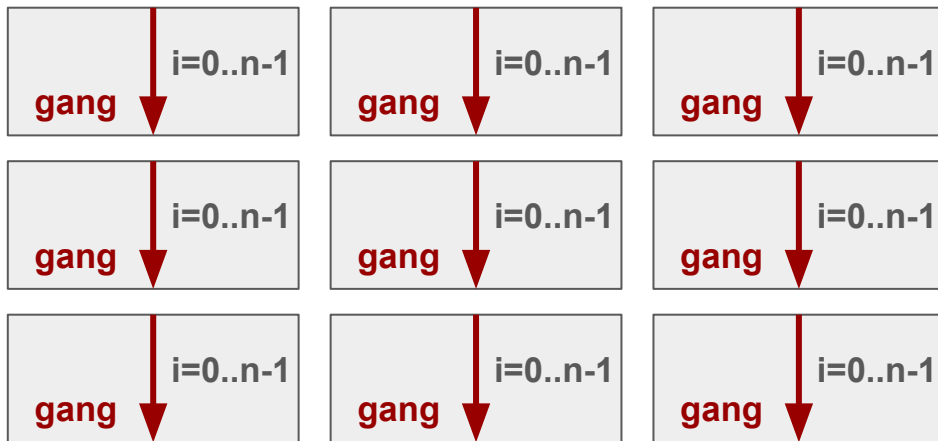
This is *not* what we want: the code is run **redundantly** we are wasting GPU resources

The parallel directive

```
!$acc parallel
    do i=1,N
        <code>
    end do
!$acc end parallel
```

The **parallel** directive:

- Creates a region where the code is run on the GPU **sequentially** (single thread)
- A **gang** is an **independent group of threads** (more details later...)



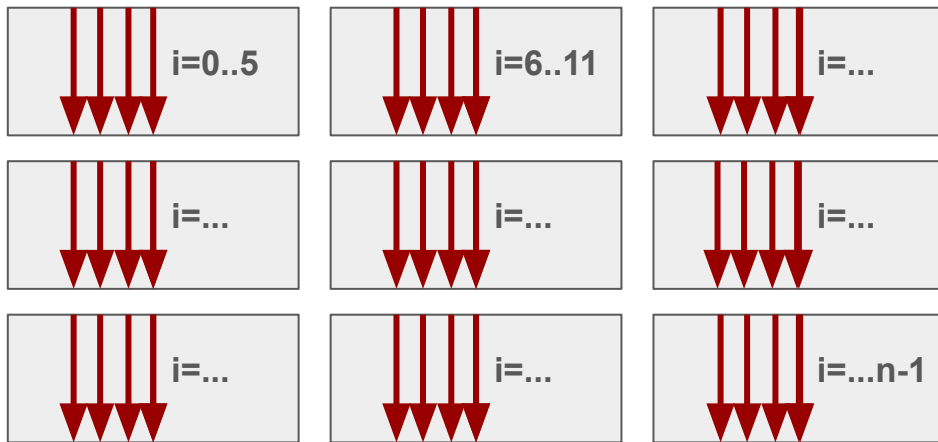
This is *not* what we want: the code is run **redundantly** we are wasting GPU resources

The loop directive

```
#pragma acc parallel
{
    #pragma acc loop
    for(i=0;i<n;i++){ <code> }
}
```

The **loop** directive:

- Instructs the compiler to distribute the loop iterations among the gangs



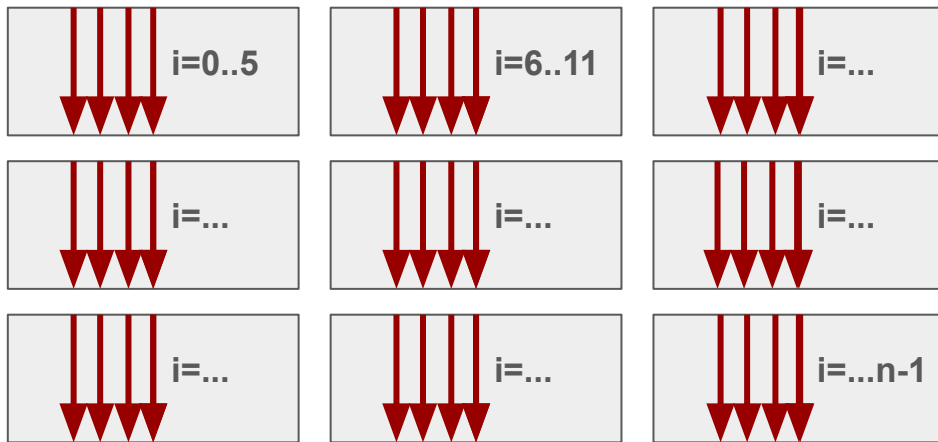
This is what we want: each gang (block) takes care of a section of the loop!

The loop directive

```
!$acc parallel
  !$acc loop
    do i=1,N
      <code>
    end do
!$acc end parallel
```

The **loop** directive:

- Instructs the compiler to distribute the loop iterations among the gangs



This is what we want: each gang (block) takes care of a section of the loop!

The parallel loop directive

```
#pragma acc parallel loop
for(i=0;i<n;i++){
    <code>
}
```

- The **parallel** and **loop** directives are often combined in the **parallel loop** directive
- The compiler might decide to insert **loop** directives to avoid race conditions!

Quick recap:

- **Parallel** marks the region of the code that will be translated into a GPU kernel (i.e: parallel execution)
- **Loop** is used to tell the compiler to parallelize the next loop it finds across the gangs

The parallel loop directive

```
!$acc parallel loop
do i=1,N
    <code>
end do
```

- The **parallel** and **loop** directives are often combined in the **parallel loop** directive
- The compiler might decide to insert **loop** directives to avoid race conditions!

Quick recap:

- **Parallel** marks the region of the code that will be translated into a GPU kernel (i.e: parallel execution)
- **Loop** is used to tell the compiler to parallelize the next loop it finds across the gangs

The kernels directive

```
#pragma acc kernels
{
    for(i = 0; i < n; i++)
        a[i] = 0;

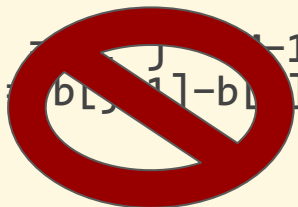
    for(j = 0; j < m-1; j++)
        b[j] = b[j+1]-b[j];
}
```

- The **kernels** directive instructs the compiler to search for parallel loops in the code
- The compiler will analyze the loops and parallelize those it finds safe to do so
- The kernels directive can be applied to regions containing multiple nested loops

The kernels directive

```
#pragma acc kernels
{
    for(i = 0; i < n; i++)
        a[i] = 0;

    for(j = 0; j < m-1; j++)
        b[j] = b[j+1]-b[j];
}
```



C/C++ allows aliasing, so you need to tell the compiler that pointers are not aliased!
(using restrict/__restrict__)

- The **kernels** directive instructs the compiler to search for parallel loops in the code
- The compiler will analyze the loops and parallelize those it finds safe to do so
- The kernels directive can be applied to regions containing multiple nested loops

The kernels directive

```
!$acc kernels  
a(:) = 1  
b(:) = 2  
c(:) = a(:) + b(:)  
!$acc end kernels
```

In Fortran aliasing is
forbidden by the standard

- The **kernels** directive instructs the compiler to search for parallel loops in the code
- The compiler will analyze the loops and parallelize those it finds safe to do so
- The kernels directive can be applied to regions containing multiple nested loops
- Compatible with Fortran array notation

Nested loops

```
#pragma acc parallel loop
for(i=0;i<n;i++){
    #pragma acc loop
    for(j=0;j<m;j++){
        a[i*n+j]=0;
    }
}
```

```
#pragma acc kernels
{
    for(i=0;i<n;i++){
        for(j=0;j<m;j++){
            a[i*n+j]=0;
        }
    }
}
```

- The **parallel loop / kernels** directives can be nested to parallelize multi-dimensional loops
- If resources are available, the compiler can exploit more level of parallelism

Nested loops

```
!$ acc parallel loop
do i=1,n
  !$ acc loop
  do j=1,m
    a(i,j)=0
  end do
end do
```

```
!$acc kernels
do i=1,n
  do j=1,m
    a(i,j)=0
  end do
end do
!$acc end kernels
```

- The **parallel loop / kernels** directives can be nested to parallelize multi-dimensional loops
- If resources are available, the compiler can exploit more level of parallelism

Nested loops: collapse clause

```
#pragma acc parallel loop collapse(2)
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        a[i*n+j]=0;
    }
}
```

```
for(ij=0;ij<n*n;ij++){
    int i = ij / n;
    int j = ij % n;
    a[i*n+j]=0;
}
}
```

- The **collapse(n)** clause collapses the next n **tightly** nested loop in a single one
- Useful for creating larger loops in order to **increase memory locality** and **exploit more parallelism**

Privatizations

```
double tmp[3];  
  
#pragma acc parallel loop private(tmp[0:3])  
for( i = 0; i < size; i++ ){  
    tmp[0] = i;  
    tmp[1] = i+1;  
    tmp[2] = i+2;  
}
```

(first)private clause allows each thread to have private copies of variables:

- **private** variables are uninitialized.
- **firstprivate** private values are initialized to the last value used on the host.

Privatizations

```
double tmp[3];
tmp[0]=1; tmp[1]=2; tmp[3]=3;
#pragma acc parallel loop private(tmp[0:3])
for( i = 0; i < n; i++ ) {
    tmp[0] = i;
    tmp[1] = i+1;
    tmp[2] = i+2;
    #pragma acc loop
    for ( j = 0; j < n; j++ ) {
        M[n*i+j] = tmp[0]+tmp[1]+tmp[2]+j;
    }
}
printf(tmp);
```

OUTPUT:
(1,2,3)

The tmp[] array is private in each iteration of the outer loop

This private copy is shared between the threads in the inner loop (same scope)

The value on the host won't be updated!

Privatizations (scalars)

- Scalars are by default private in kernel regions and firstprivate in parallel regions
- You typically don't have to touch scalars except (e.g.):
 - 1.) Global vars in C/C++, module vars in Fortran
 - 2.) Scalar used as an rvalue after the parallel region

The reduction clause

```
#pragma acc parallel loop
for (int i = 0; i < N; i++) {
    dot += a[i] * b[i];
}
```

The reduction clause

```
#pragma acc parallel loop  
for (int i = 0; i < N; i++) {  
    dot += a[i] * b[i];  
}
```

Scalars are
firstprivate: the
value of dot is
not changed on
the host!

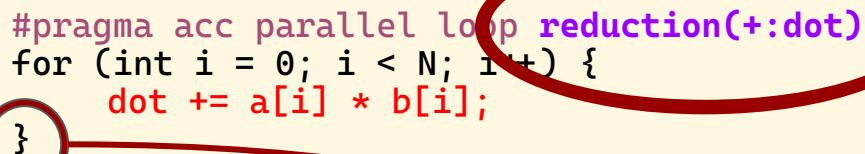
The reduction clause

```
#pragma acc parallel loop reduction(+:dot)
for (int i = 0; i < N; i++) {
    dot += a[i] * b[i];
}
```

- Each thread will have its own **private copy** and will perform a **partial reduction**

The reduction clause

```
#pragma acc parallel loop reduction(+:dot)
for (int i = 0; i < N; i++) {
    dot += a[i] * b[i];
}
```



- Each thread will have its own **private copy** and will perform a **partial reduction**

- At the end, the local sums are combined safely

The compiler is very keen on inserting reduction clauses where needed, e.g. in this case it would have done so automatically.

The reduction clause

```
#pragma acc parallel loop reduction(+:v[0])  
for (int i = 0; i < N; i++) {  
    v[0] += a[i] * b[i];  
}
```

```
double tmp = v[0];  
#pragma acc parallel loop reduction(+:tmp)  
for (int i = 0; i < N; i++) {  
    tmp += a[i] * b[i];  
}  
v[0] = tmp;
```

You cannot reduce on

- Array entries
- C/C++ class/struct members
- Fortran derived types members

The reduction clause

```
#pragma acc parallel loop reduction(+:mesh.data[0])
for (int i = 0; i < N; i++) {
    mesh.data[0] += a[i] * b[i];
}
```

```
auto tmp = mesh.data[0];
#pragma acc parallel loop reduction(+:tmp)
for (int i = 0; i < N; i++) {
    tmp += a[i] * b[i];
}
mesh.data[0] = tmp;
```

You cannot reduce on

- Array entries
- C/C++ class/struct members
- Fortran derived types members

The reduction clause

```
!$acc parallel loop reduction(+:mytype%member)
do i = 1, N
    mytype%member = mytype%member + 7 + a(i)
end do
```

```
tmp = mytype%member
!$acc parallel loop reduction(+:tmp)
do i = 1, N
    tmp = tmp + 7 + a(i)
end do
mytype%member = tmp
```

You cannot reduce on

- Array entries
- C/C++ class/struct members
- Fortran derived types members

The reduction clause

Operator	Description	Example
+	Addition/Summation	<code>reduction(+:sum)</code>
*	Multiplication/Product	<code>reduction(*:product)</code>
max	Maximum value	<code>reduction(max:maximum)</code>
min	Minimum value	<code>reduction(min:minimum)</code>
&	Bitwise and	<code>reduction(&:val)</code>
 	Bitwise or	<code>reduction(:val)</code>
&&	Logical and	<code>reduction(&&:val)</code>
 	Logical or	<code>reduction(:val)</code>

Seq and atomic

```
#pragma acc parallel loop seq
for (int i = 0; i < N; i++) {
    <code>
}
```

The **seq** clause will tell the compiler to run the loop sequentially

!The compiler may insert **seq** clause (check Minfo)

```
#pragma acc parallel loop
for (int i = 0; i < N; i++) {
    if(a[i] > thr){
        #pragma acc atomic
        n_thr++
    }
}
```

The **atomic** directive will tell the compiler that the threads have to execute the instruction **sequentially**

With **atomic** all threads update the same memory location, no partial accumulation as in **reduction**

Seq and atomic

```
#pragma acc parallel loop seq
for (int i = 0; i < N; i++) {
    <code>
}
```

The **seq** clause will tell the compiler to run the loop sequentially

!The compiler may insert **seq** clause (check Minfo)

```
#pragma acc parallel loop
for (int i = 0; i < N; i++) {
    if(a[i] > thr){
        #pragma acc atomic update
        n_thr++
    }
}
```

atomic supports several clauses:

- **read/write** value=x
- **update (default)** x= expr
- **capture** x= expr; value=x

Atomic example

```
//count[i] are some integers
```

```
#pragma acc parallel loop  
for(int i=0;i<range;i++)  
    hist[i]=0;
```

```
#pragma acc parallel loop  
for(int i=0;i<d;i++) {  
    #pragma acc atomic  
        hist[count[i]]+=1;  
}
```

```
//count(i) are some integers
```

```
!$acc parallel loop  
do i=1,range  
    hist(i)=0  
end do
```

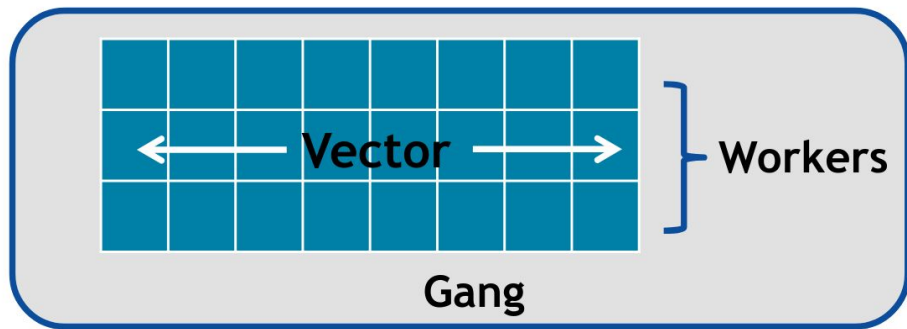
```
!$acc parallel loop  
do i=1,d  
    !$acc atomic  
        hist(count(i))+=1;  
end do
```

Layers of parallelism

OpenACC expresses 3 layers of parallelism: **gangs,workers,vectors**

- **Vector:** the finest grain
- **Gang:** the coarsest level
- **Worker:** intermediate level
- The correspondence between the openACC expressed parallelism and the hardware parallelism depends on the hardware (e.g. on Intel multicore: a gang is a thread and the vector is an AVX vector instruction)

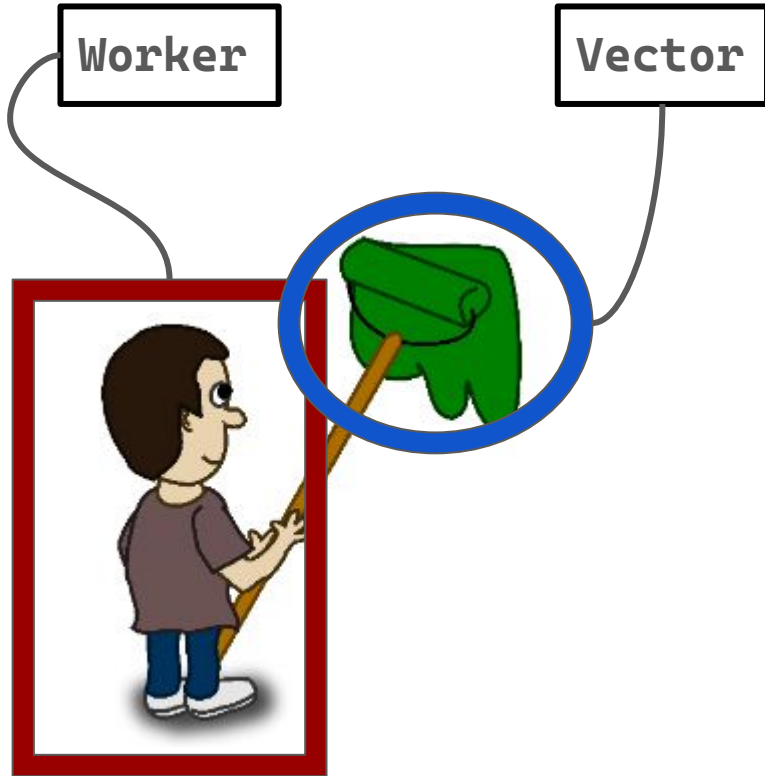
Layers of parallelism




- **Vector** : ThreadIdx.x
- **Worker** : ThreadIdx.y
- **Gang** : Block

- The size of a gang is (`num_workers * vector_lenght`)
- `vector_lenght` must be multiple of 32 (warp size)
- `These parameters` can be specified in the code as clauses

Layers of parallelism



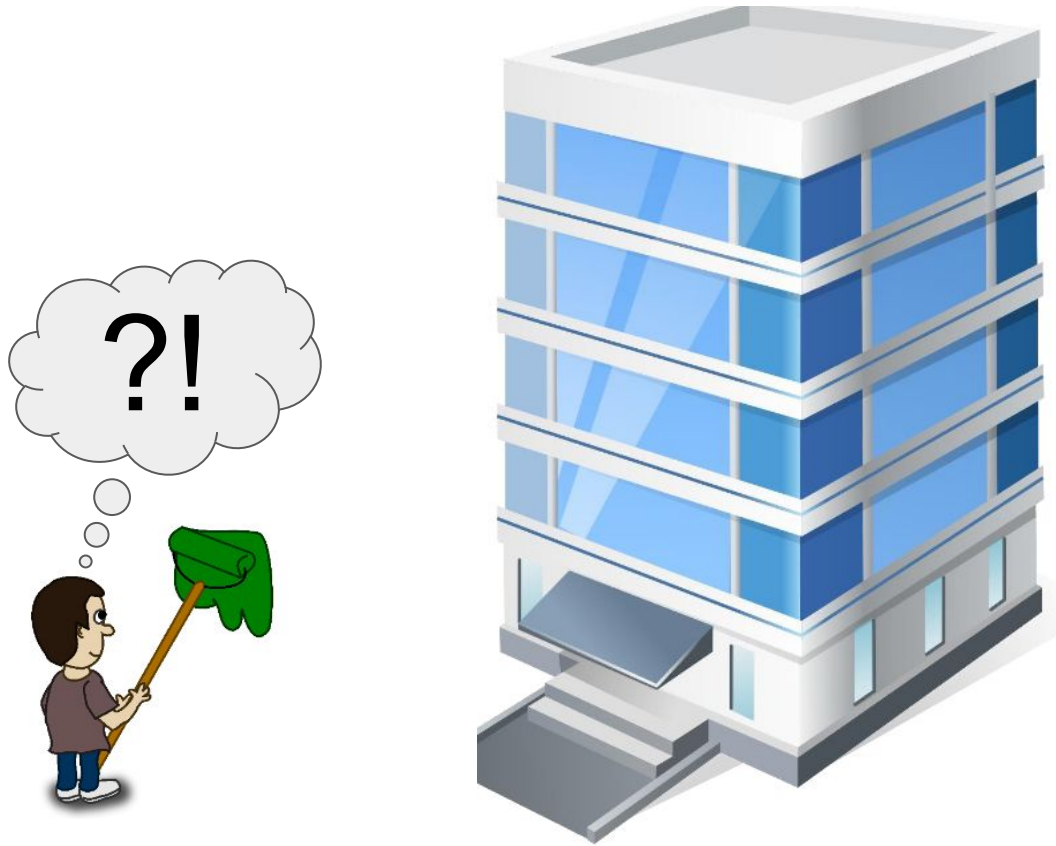
Layers of parallelism



Layers of parallelism




Layers of parallelism



Layers of parallelism



Layers of parallelism



Layers of parallelism

- Independent groups of workers (**gangs**) can clean different floors at the same time



Specify the parallelism

parallel:

- num_gangs(#)
- num_workers(#)
- vector_length(#)

kernel:

- gang(#)
- worker(#)
- vector(#)

```
#pragma acc parallel loop gang
for( i = 0; i < N; i++ )
    #pragma acc loop vector
    for( j = 0; j < M; j++ )
        < loop code >
```

- Typically the compiler generates G gangs, one worker and a vector size of 128

Specify the parallelism

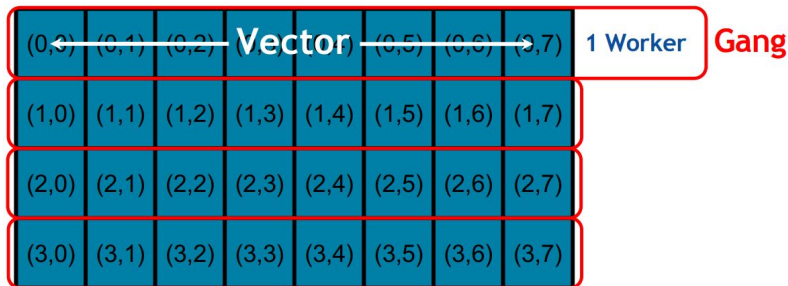
parallel:

- **num_gangs(#)**
- **num_workers(#)**
- **vector_length(#)**

kernel:

- **gang(#)**
- **worker(#)**
- **vector(#)**

```
#pragma acc parallel loop gang
for( i = 0; i < N; i++ )
    #pragma acc loop vector
    for( j = 0; j < M; j++ )
        < loop code >
```



Specify the parallelism

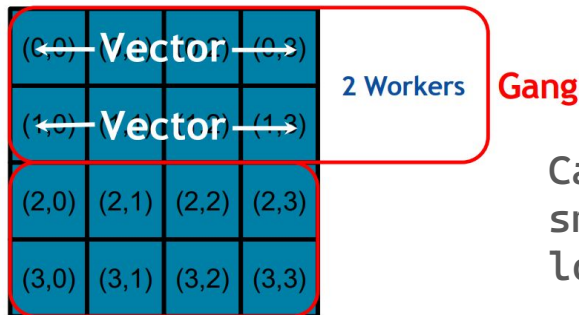
parallel:

- `num_gangs(#)`
- `num_workers(#)`
- `vector_length(#)`

kernel:

- `gang(#)`
- `worker(#)`
- `vector(#)`

```
#pragma acc parallel num_workers(2)
#pragma acc loop gang worker
for( i = 0; i < N; i++ )
    #pragma acc loop vector
    for( j = 0; j < M; j++ )
        < loop code >
```



Can be useful for
smaller inner
loops

THANK YOU

FOR YOUR ATTENTION!

Exercises

1. Run the examples and play around with the pragmas
2. Implement pi estimation with openACC
3. Implement matrix multiplication using openACC pragmas (try different ways)
4. (Advanced weekend exercise ^_^) Implement mat mul with openACC using different configurations of gangs,worker,vector and compare performance